

Artificial Intelligence: Representation and Problem Solving

15-381


January 16, 2007

Neural Networks

Topics

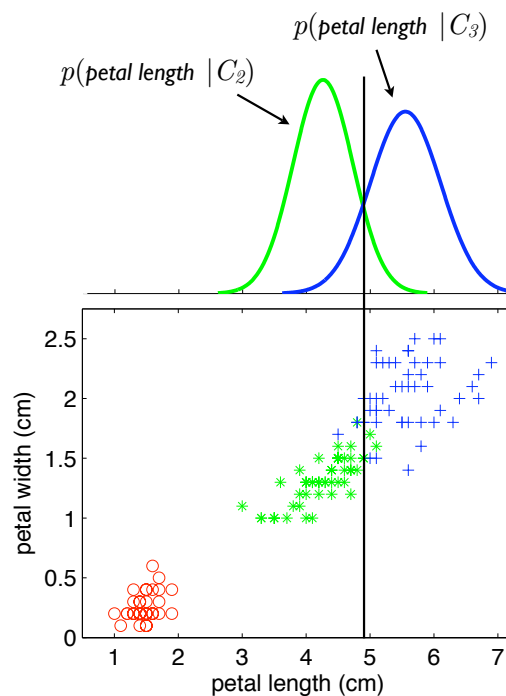
- decision boundaries
- linear discriminants
- perceptron
- gradient learning
- neural networks

The Iris dataset with decision tree boundaries




The optimal decision boundary for C_2 vs C_3

- optimal decision boundary is determined from the statistical distribution of the classes
- optimal only if model is correct
- assigns precise degree of uncertainty to classification




Optimal decision boundary




Can we do better?

- only way is to use more information
- DTs use both petal width and petal length



Arbitrary decision boundaries would be more powerful




Defining a decision boundary

- consider just two classes
- want points on one side of line in class 1, otherwise class 2.
- 2D linear discriminant function:

$$\begin{aligned} y &= \mathbf{m}^T \mathbf{x} + b \\ &= m_1 x_1 + m_2 x_2 + b \\ &= \sum_i m_i x_i + b \end{aligned}$$

- This defines a 2D plane which leads to the decision:

$$\mathbf{x} \in \begin{cases} \text{class 1} & \text{if } y \geq 0, \\ \text{class 2} & \text{if } y < 0. \end{cases}$$



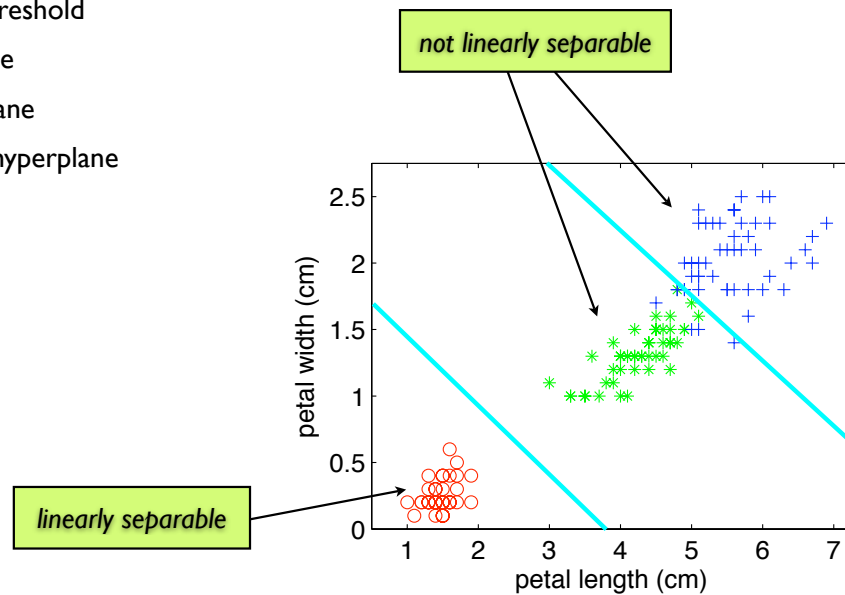
The decision boundary:
 $y = \mathbf{m}^T \mathbf{x} + b = 0$

Or in terms of scalars:

$$\begin{aligned} m_1 x_1 + m_2 x_2 &= -b \\ \Rightarrow x_2 &= -\frac{m_1 x_1 + b}{m_2} \end{aligned}$$

Linear separability

- Two classes are linearly separable if they can be separated by a linear combination of attributes
 - 1D: threshold
 - 2D: line
 - 3D: plane
 - M-D: hyperplane




Diagramming the classifier as a “neural” network

- The feedforward neural network is specified by weights w_i and bias b :

$$y = \mathbf{w}^T \mathbf{x} + b$$


$$= \sum_{i=1}^M w_i x_i + b$$



- It can be written equivalently as

$$y = \mathbf{w}^T \mathbf{x} = \sum_{i=0}^M w_i x_i$$

- where $w_0 = b$ is the bias and a “dummy” input x_0 that is always 1.



Determining, ie learning, the optimal linear discriminant

- First we must define an *objective function*, ie the goal of learning
- Simple idea: adjust weights so that output $y(\mathbf{x}_n)$ matches class c_n
- *Objective*: minimize sum-squared error over all patterns \mathbf{x}_n :

$$E = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n)^2$$

- Note the notation \mathbf{x}_n defines a *pattern vector*:


$$\mathbf{x}_n = \{x_1, \dots, x_M\}_n$$

- We can define the desired class as:

$$c_n = \begin{cases} 0 & \mathbf{x}_n \in \text{class 1} \\ 1 & \mathbf{x}_n \in \text{class 2} \end{cases}$$

We've seen this before: curve fitting

$$t = \sin(2\pi x) + \text{noise}$$

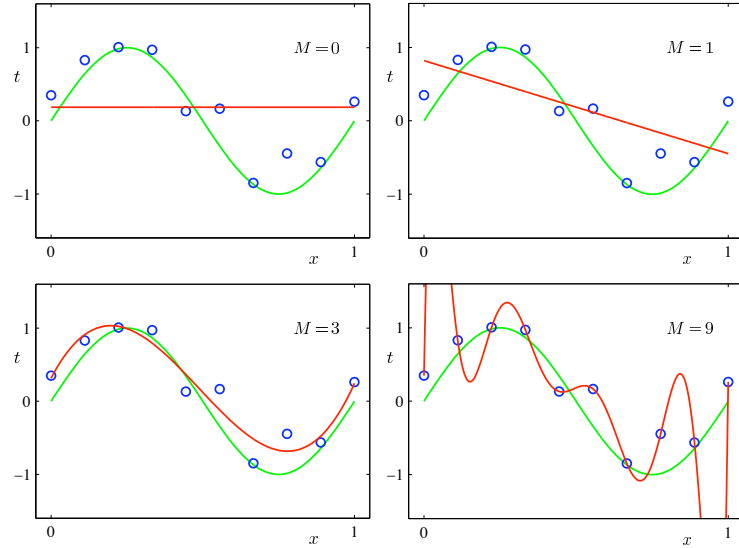


Neural networks compared to polynomial curve fitting

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y(x_n, \mathbf{w}) - t_n]^2$$

For the linear network, $M=1$ and there are multiple input dimensions



example from Bishop (2006), *Pattern Recognition and Machine Learning*

General form of a linear network

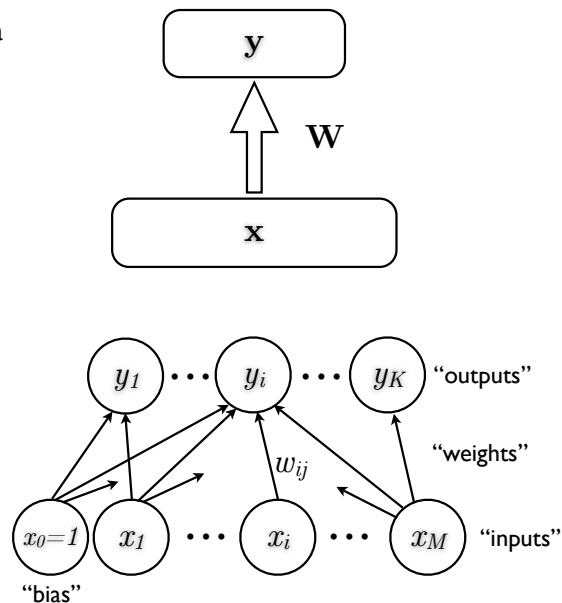
- A linear neural network is simply a linear transformation of the input.

$$y_j = \sum_{i=0}^M w_{i,j}x_i$$

- Or, in matrix-vector form:

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

- Multiple outputs corresponds to *multivariate regression*



Training the network: Optimization by gradient descent

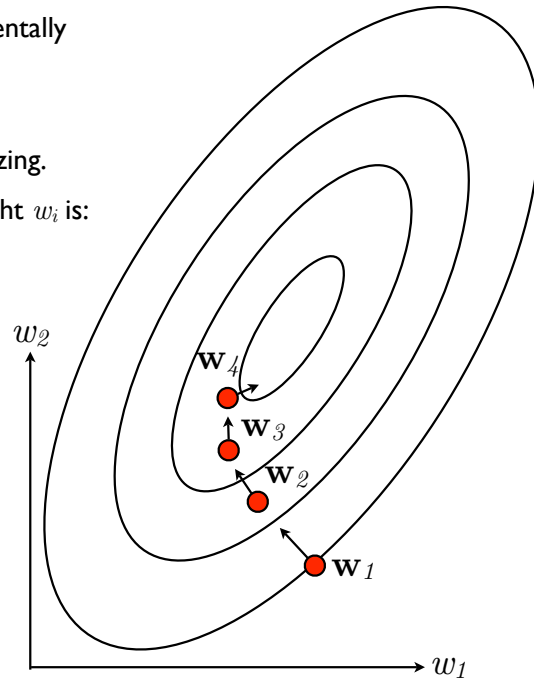
- We can adjust the weights incrementally to minimize the objective function.
- This is called *gradient descent*
- Or *gradient ascent* if we're maximizing.
- The gradient descent rule for weight w_i is:

$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

- Or in vector form:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \epsilon \frac{\partial E}{\partial \mathbf{w}}$$

- For *gradient ascent*, the *sign* of the gradient step changes.



Computing the gradient

- Idea: minimize error by gradient descent
- Take the derivative of the objective function wrt the weights:

$$E = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n)^2$$

$$\frac{\partial E}{\partial w_i} = \frac{2}{2} \sum_{n=1}^N (w_0 x_{0,n} + \dots + w_i x_{i,n} + \dots + w_M x_{M,n} - c_n) x_{i,n}$$

$$= \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- And in vector form:

$$\frac{\partial E}{\partial \mathbf{w}} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) \mathbf{x}_n$$

Simulation: learning the decision boundary

- Each iteration updates the gradient:


$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- Epsilon is a small value:

$$\epsilon = 0.1/N$$

- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary

- Each iteration updates the gradient:

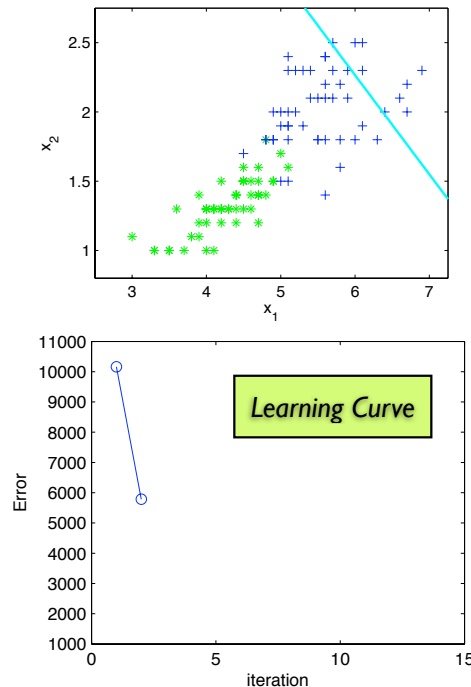
$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- Epsilon is a small value:

$$\epsilon = 0.1/N$$

- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary

- Each iteration updates the gradient:



$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- Epsilon is a small value:

$$\epsilon = 0.1/N$$

- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary

- Each iteration updates the gradient:



$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- Epsilon is a small value:

$$\epsilon = 0.1/N$$

- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary

- Each iteration updates the gradient:

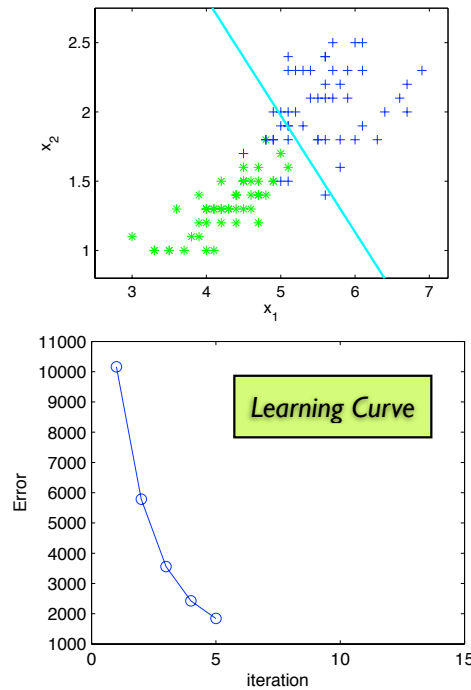
$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$

- Epsilon is a small value:

$$\epsilon = 0.1/N$$

- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary

- Each iteration updates the gradient:

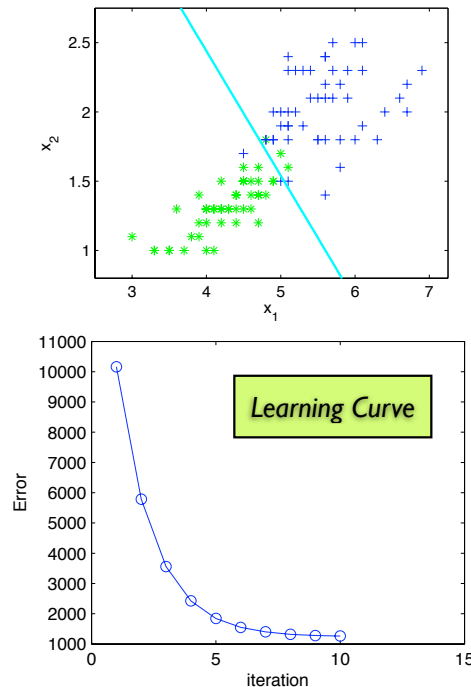
$$w_i^{t+1} = w_i^t - \epsilon \frac{\partial E}{\partial w_i}$$

$$\frac{\partial E}{\partial w_i} = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - c_n) x_{i,n}$$


- Epsilon is a small value:

$$\epsilon = 0.1/N$$


- Epsilon too large:
 - learning diverges
- Epsilon too small:
 - convergence slow



Simulation: learning the decision boundary




- Learning converges onto the solution that minimizes the error.
- For linear networks, this is guaranteed to converge to the minimum
- It is also possible to derive a closed-form solution (covered later)




Learning is slow when epsilon is too small

- Here, larger step sizes would converge more quickly to the minimum



Divergence when epsilon is too large

- If the step size is too large, learning can oscillate between different sides of the minimum



Multi-layer networks

- Can we extend our network to multiple layers? We have:


$$y_j = \sum_i w_{i,j} x_i$$

$$\begin{aligned} z_j &= \sum_k v_{j,k} y_k \\ &= \sum_k v_{j,k} \sum_i w_{i,k} x_i \end{aligned}$$

- Or in matrix form

$$\begin{aligned} \mathbf{z} &= \mathbf{V}\mathbf{y} \\ &= \mathbf{V}\mathbf{W}\mathbf{x} \end{aligned}$$

- Thus a two-layer linear network is equivalent to a one-layer linear network with weights $\mathbf{U} = \mathbf{V}\mathbf{W}$.
- It is *not* more powerful.



How do we address this?

Non-linear neural networks

- Idea introduce a non-linearity:



$$y_j = f\left(\sum_i w_{i,j} x_i\right)$$

- Now, multiple layers are *not* equivalent


$$\begin{aligned} z_j &= f\left(\sum_k v_{j,k} y_k\right) \\ &= f\left(\sum_k v_{j,k} f\left(\sum_i w_{i,k} x_i\right)\right) \end{aligned}$$

- Common nonlinearities:


- threshold
- sigmoid




Modeling logical operators



x_1 AND x_2




x_1 OR x_2



x_1 XOR x_2


- A one-layer binary-threshold network can implement the logical operators AND and OR, but not XOR.
- Why not?

$$y_j = f\left(\sum_i w_{i,j} x_i\right)$$



Posterior odds interpretation of a sigmoid

The general classification/regression problem



Given data, we want to learn a model that can correctly classify novel observations **or** **map the inputs to the outputs**

A general multi-layer neural network

- Error function is defined as before, where we use the target vector \mathbf{t}_n to define the desired output for network output \mathbf{y}_n .

$$E = \frac{1}{2} \sum_{n=1}^N (\mathbf{y}_n(\mathbf{x}_n, \mathbf{W}_{1:L}) - \mathbf{t}_n)^2$$

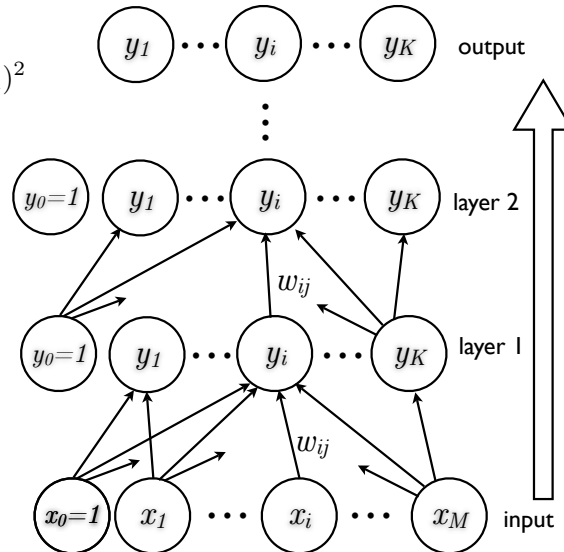
- The “forward pass” computes the outputs at each layer:

$$y_j^l = f\left(\sum_i w_{i,j}^l y_i^{l-1}\right)$$

$$l = \{1, \dots, L\}$$

$$\mathbf{x} \equiv \mathbf{y}^0$$

$$\text{output} = \mathbf{y}^L$$



Deriving the gradient for a sigmoid neural network

- Mathematical procedure for train is gradient descent: same as before, except the gradients are more complex to derive.

New problem: local minima

$$E = \frac{1}{2} \sum_{n=1}^N (\mathbf{y}_n(\mathbf{x}_n, \mathbf{W}_{1:L}) - \mathbf{t}_n)^2$$

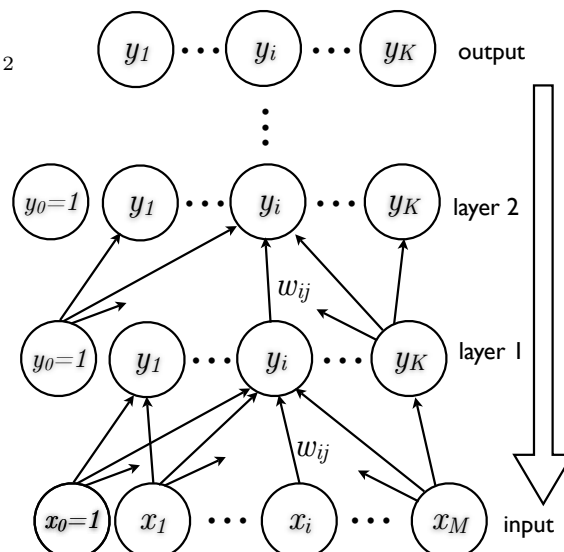
- Convenient fact for the sigmoid non-linearity:

$$\frac{d\sigma(x)}{dx} = \frac{d}{dx} \frac{1}{1 + \exp(-x)}$$

$$= \sigma(x)(1 - \sigma(x))$$


- backward pass computes the gradients: *back-propagation*

$$\mathbf{W}_{t+1} = \mathbf{W}_t + \epsilon \frac{\partial E}{\partial \mathbf{W}}$$




Applications: Driving (output is analog: steering direction)

network with 1 layer
(4 hidden units)



- Learns to drive on roads
- Demonstrated at highway speeds over 100s of miles




D. Pomerleau. *Neural network perception for mobile robot guidance*. Kluwer Academic Publishing, 1993.

Real image input is augmented to avoid overfitting


Training data:
Images +
corresponding
steering angle



Important:
Conditioning of
training data to
generate new
examples → avoids
overfitting




Hand-written digits: LeNet



- Takes as input image of handwritten digit
- Each pixel is an input unit
- Complex network with many layers
- Output is digit class
- Tested on large (50,000+) database of handwritten samples
- Real-time
- Used commercially

LeNet




Very low error rate (<< 1%)

Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, november 1998.

<http://yann.lecun.com/exdb/lenet/>

Object recognition



- LeCun, Huang, Bottou (2004). Learning Methods for Generic Object Recognition with Invariance to Pose and Lighting. Proceedings of CVPR 2004.
- <http://www.cs.nyu.edu/~yann/research/norb/>

Summary

- Decision boundaries
 - Bayes optimal
 - linear discriminant
 - linear separability
- Classification vs regression
- Optimization by gradient descent
- Degeneracy of a multi-layer linear network
- Non-linearities: threshold, sigmoid, others?
- Issues:
 - very general architecture, can solve many problems
 - large number of parameters: need to avoid overfitting
 - usually requires a large amount of data, or special architecture
 - local minima, training can be slow, need to set stepsize